Names: Yazmyn Sims, Annie McKay

Network Training

● a printout showing the shape of your network, as generated by model.print_summary() marked “Initial Network”

```
* Training basic_model for 5 epochs
Model: "sequential"
_________________________________________________________________
 Layer (type)                Output Shape              Param #
=================================================================
 rescaling (Rescaling)       (None, 150, 150, 3)       0

 conv2d (Conv2D)             (None, 148, 148, 8)       224

 max_pooling2d (MaxPooling2D  (None, 74, 74, 8)        0
 )

 conv2d_1 (Conv2D)           (None, 72, 72, 8)         584

 max_pooling2d_1 (MaxPooling  (None, 36, 36, 8)        0
 2D)

 conv2d_2 (Conv2D)           (None, 34, 34, 8)         584

 max_pooling2d_2 (MaxPooling  (None, 17, 17, 8)        0
 2D)

 flatten (Flatten)           (None, 2312)              0

 dense (Dense)               (None, 64)                148032

 dropout (Dropout)           (None, 64)                0

 dense_1 (Dense)             (None, 3)                 195

=================================================================
Total params: 149,619
Trainable params: 149,619
Non-trainable params: 0
```

- **a plot showing training and validation loss as function of epoch**

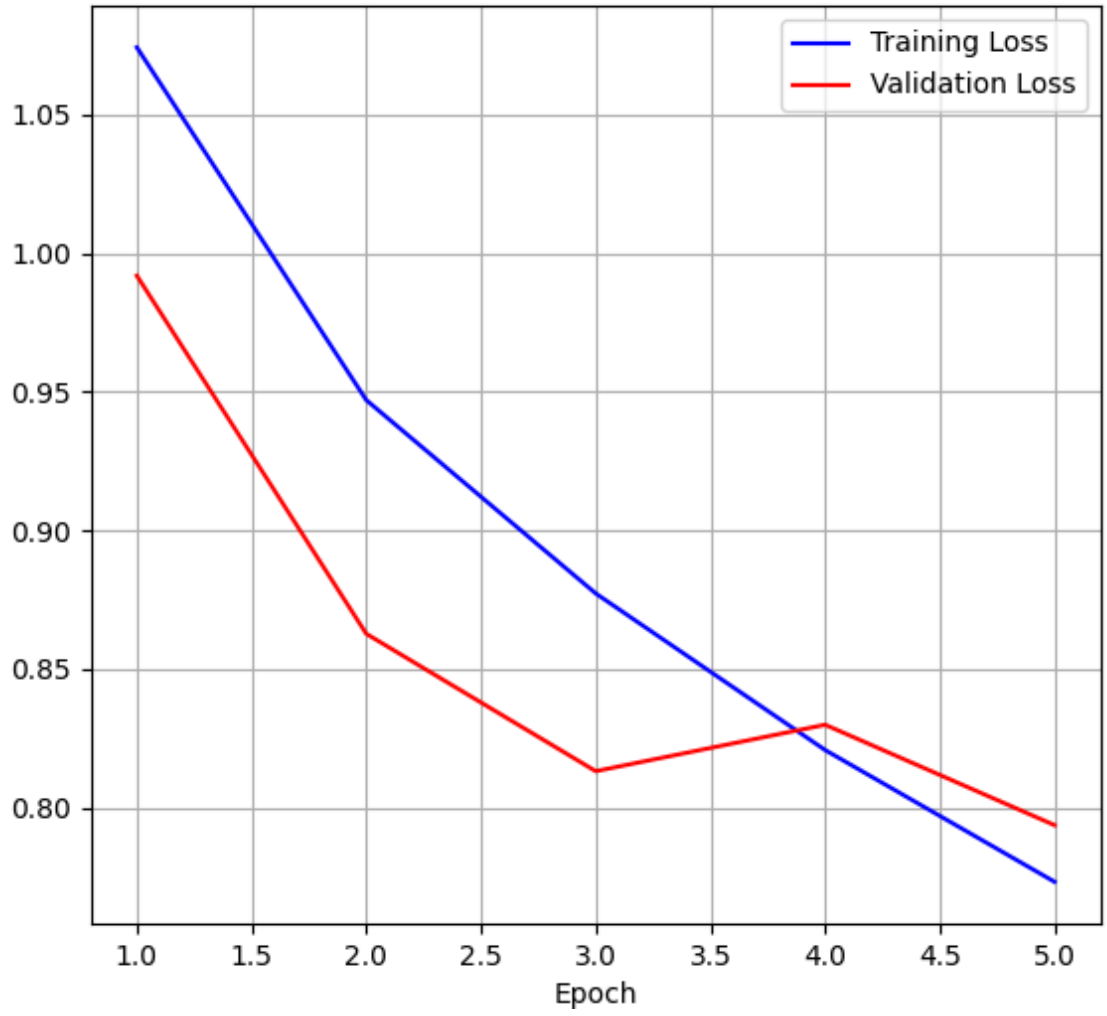

● **a plot showing accuracy against the training and validation sets as a function of epoch**

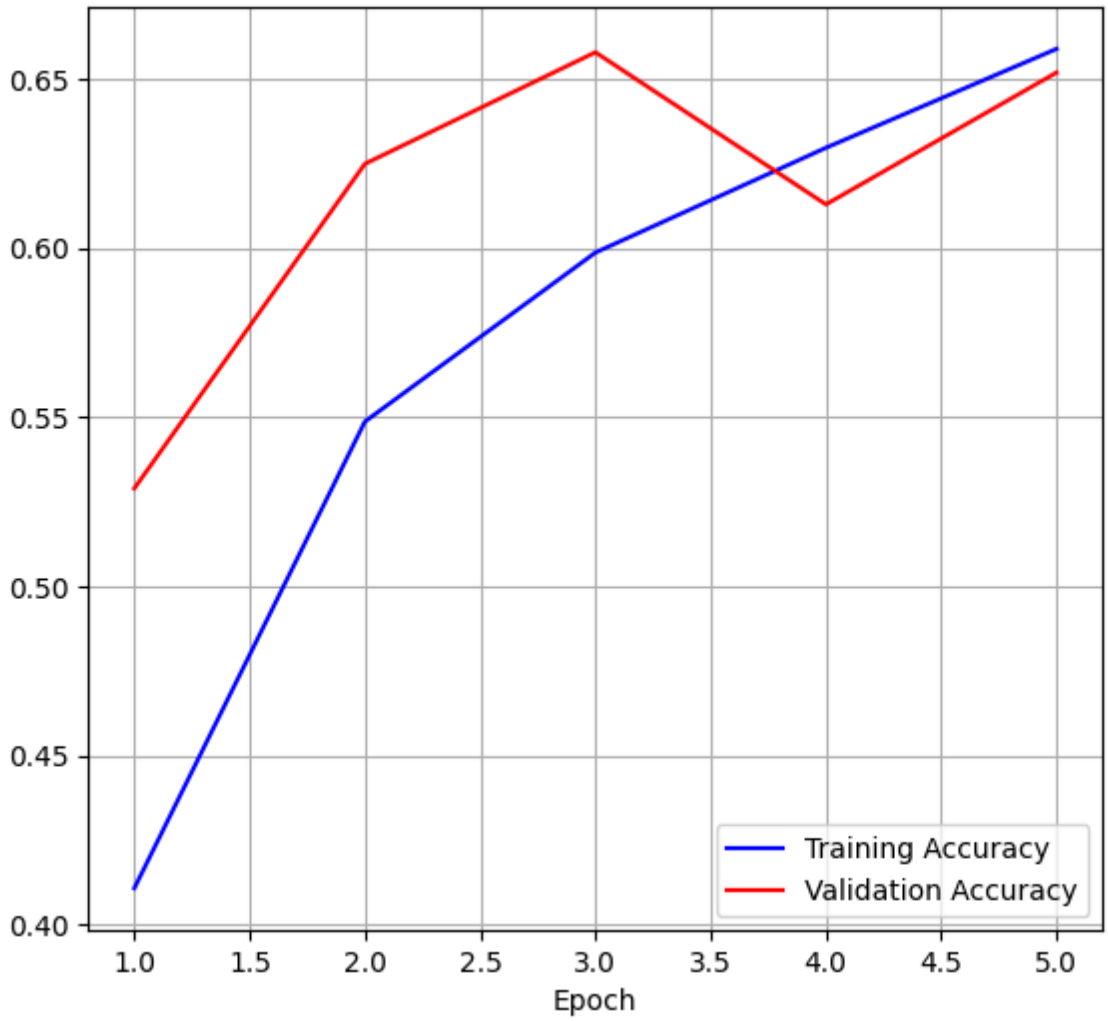

● **the accuracy and loss of your best learned model (obtained as the model in effect when overfitting begins) when measured against the held-back test set**. Our best learned model is on Epoch 3 right before overfitting begins. It has a 66% Training accuracy and a 60% validation accuracy. At the same time the loss for both training and validation is under 0.9%